

Implementing Clustering and Dimensionality Reduction in scikit-learn



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Clustering is an unsupervised learning technique which helps find patterns in data

Common clustering algorithms are k-means, mean-shift clustering

Dimensionality reduction represents inputs in terms of their most significant features

PCA is a very commonly used technique for latent factor analysis

Types of ML Algorithms



Supervised

Labels associated with the training data is used to correct the algorithm



Unsupervised

The model has to be set up right to learn structure in the data

Types of ML Algorithms



Supervised

Labels associated with the training data is used to correct the algorithm



Unsupervised

The model has to be set up right to learn structure in the data

Clustering

Clustering




Anything
can be
represented
by a set of
numbers

Age, Height, Weight




35




Age



Age, Height, Weight





A set of N numbers represents
a point in an **N-dimensional**
Hypercube

Clustering



**A set of points, each
representing a Facebook user**

Clustering



Same group = **similar**

Different group = **different**

Clustering



Same group = **similar**



Different group = **different**

Clustering




The **distance** between users
in a cluster indicates how
similar they are

Clustering




Maximize **intra**-cluster
similarity

Clustering



Minimize **inter**-cluster
similarity

Clustering Objective




Maximize intra-cluster similarity

Minimize inter-cluster similarity

The **K-Means Clustering** algorithm
is a famous Machine Learning
algorithm to achieve this


K-Means Clustering

**Initialize K
centroids i.e.
means**




K-Means Clustering

Assign
each point
to a cluster




K-Means Clustering

Assign
each point
to a cluster




K-Means Clustering

Assign
each point
to a cluster




K-Means Clustering

Assign
each point
to a cluster




**Recalculate
the mean
for each
cluster**

K-Means Clustering




K-Means Clustering

**Re-assign
the points
to clusters**




**Iterate until
points are
in their final
clusters**

K-Means Clustering




K-Means Clustering




Convergence

K-Means Clustering



K-Means Clustering



Each cluster has a representative
point called a **reference vector**


K-Means Clustering



Because of how they are
calculated, these reference
vectors are often called **centroids**


Mean Shift Clustering

**Start with a
set of points
in space**




Mean Shift Clustering

**Define a
neighborhood
for each point**




Mean Shift Clustering

**Define a
neighborhood
for each point**



Mean Shift Clustering


**Define a
neighborhood
for each point**



Mean Shift Clustering

For each point, calculate a function based on all points in the neighborhood


That function is called the **kernel**



Flat Kernel


Flat kernel: sum of all
points in neighborhood

Each point gets the
same weight




Gaussian (RBF) Kernel

**Probability-weighted
sum of points**



**What probability
distribution?**

Gaussian (RBF) Kernel




Gaussian probability distribution

Defined by


- mean μ
- standard deviation σ

Gaussian Distribution



$$N(\mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Mean = Center point



Gaussian (RBF) Kernel


Mean μ = center point

Standard deviation $\sigma \sim$ bandwidth

(Bandwidth is a hyperparameter)


Mean Shift Clustering

**Kernel is
applied to
each point**




Mean Shift Clustering

**Kernel is
applied to
each point**




Mean Shift Clustering

**Kernel is
applied to
each point**




Mean Shift Clustering

**Kernel is
applied to
each point**




Mean Shift Clustering

Assume points are
color-coded by
magnitude of RBF




Mean Shift Clustering

**High RBF
values are
peaks**




Mean Shift Clustering

**Low RBF
values are
troughs**




Mean Shift Clustering

Now, all points start
to “shift” towards
the nearest peak




Mean Shift Clustering

Now, all points start
to “shift” towards
the nearest peak




Mean Shift Clustering

Now, all points start
to “shift” towards
the nearest peak




Mean Shift Clustering

**This is the
“mean shift”**




Mean Shift Clustering

**This is the
“mean shift”**




Mean Shift Clustering

**Algorithm
converges when
points stop moving**



Role of Bandwidth



$$\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$


Standard deviation $\sigma \sim$ bandwidth

Bandwidth is the only hyperparameter

Small bandwidth \sim tall skinny kernel

Large bandwidth \sim flat kernel

Role of Bandwidth



Tall skinny kernel

Ignore points far from the mean



Flatter kernel

Considers points far from the mean

Similar, yet Different

K-Means Clustering

Need to specify number of clusters as hyperparameter

Can't handle some complex non-linear data

Less hyperparameter tuning needed

Mean Shift Clustering

No need to specify number of clusters upfront as hyperparameter

Uses density function to handle even complex non-linear data (e.g. pixels)

Hyperparameter tuning very important

Similar, yet Different

K-Means Clustering

Computationally less intensive

$O(N)$ in number of data points

Struggles with outliers

Mean Shift Clustering

Computationally very intensive

$O(N^2)$ in number of data points

Copes better with outliers

Demo

**Implement mean-shift clustering in
scikit-learn**

Principal Components Analysis

Principal Components Analysis


A technique to re-express complex data in terms of a few, well-chosen vectors (Principal Components) that most efficiently capture the variation in that data

Intuition Behind PCA




Objective: Find the “best” directions to represent this data

Intuition Behind PCA




Start by “projecting” the data onto a line in some direction

Intuition Behind PCA




Start by “projecting” the data onto a line in some direction

Intuition Behind PCA



The greater the distances between these projections,
the “better” the direction

Bad Projection




A projection where the distances are minimised is a bad one - **information is lost**

Good Projection




A projection where the distances are maximised is a good one - **information is preserved**

Intuition Behind PCA




The direction along which this variance is maximised is the **first principal component** of the original data

Intuition Behind PCA




Find the next best direction, the **second principal component**, which must be at right angles to the first

Intuition Behind PCA




Find the next best direction, the **second principal component**, which must be at right angles to the first

Principal Components at Right Angles




Directions at right angles help express the most variation with the smallest number of directions

Intuition Behind PCA




The variances are clearly smaller along this **second principal component** than along the first

Intuition Behind PCA




In general, there are as many principal components as there are dimensions in the original data

Intuition Behind PCA




Re-orient the data along these new axes

Dimensionality Reduction




If the **variance** along the second principal component is small enough, we can just **ignore** it and use just 1 dimension to represent the data

Dimensionality Reduction



Variation along 2 dimensions: 2 principal components required

Dimensionality Reduction



Variation along 1 dimension: 1 principal component is sufficient


Principal Components Analysis

A technique to re-express complex data in terms of a few, well-chosen vectors (Principal Components) that most efficiently capture the variation in that data

Data of high dimensionality, each point represented as $(x_1, x_2 \dots x_N)$

Principal Components Analysis

A technique to re-express **complex data** in terms of a few, well-chosen vectors (Principal Components) that most efficiently capture the variation in that data



Principal Components Analysis

A technique to re-express complex data in terms of a few, well-chosen vectors (Principal Components) that most efficiently capture the variation in that data



These define a smaller number of new dimensions, e.g. just two (F_1 , F_2)

Express each original point
 $(x_1, x_2 \dots x_N)$ as just (f_1, f_2)

Principal Components Analysis

A technique to re-express complex data in terms of a few, well-chosen vectors (Principal Components) that most efficiently capture the variation in that data

Principal Components Analysis


A technique to re-express complex data in terms of a few, well-chosen vectors (Principal Components) that most efficiently capture the variation in that data



**Very little information from
the original data is lost**


Principal Components Analysis

A technique to re-express complex data in terms of a few, well-chosen vectors (Principal Components) that most **efficiently** capture the variation in that data



Principal Components are a very efficient representation of the original data

Principal Components Analysis




Original Data

Same number of
columns

Principal
Components

Dimensionality Reduction



Reconstruct Original Data

The diagram illustrates the reconstruction of original data from principal components and weight vectors. It consists of three main parts: Principal Components, Weight Vectors, and Original Data, connected by a multiplication and equality sign.

Principal Components: A matrix of size $n \times k$ containing elements $f_{11}, f_{1k}, f_{21}, f_{2k}, f_{31}, f_{3k}, \dots, f_{n1}, f_{nk}$. The width is labeled k and the height is labeled n .

Weight Vectors: A matrix of size $k \times k$ containing elements $w_{11}, w_{1k}, w_{21}, w_{2k}, w_{31}, w_{3k}, \dots, w_{kk}, w_{kk}$. The width is labeled k and the height is labeled k .

Original Data: A matrix of size $n \times k$ containing elements $x_{11}, x_{1k}, x_{21}, x_{2k}, x_{31}, x_{3k}, \dots, x_{n1}, x_{nk}$. The width is labeled k and the height is labeled n .

The equation is represented as:

$$\begin{bmatrix} f_{11} & \dots & f_{1k} \\ f_{21} & \dots & f_{2k} \\ f_{31} & \dots & f_{3k} \\ \dots & \dots & \dots \\ f_{n1} & \dots & f_{nk} \end{bmatrix} \times \begin{bmatrix} w_{11} & \dots & w_{1k} \\ w_{21} & \dots & w_{2k} \\ w_{31} & \dots & w_{3k} \\ \dots & \dots & \dots \\ w_{kk} & \dots & w_{kk} \end{bmatrix} = \begin{bmatrix} x_{11} & \dots & x_{1k} \\ x_{21} & \dots & x_{2k} \\ x_{31} & \dots & x_{3k} \\ \dots & \dots & \dots \\ x_{n1} & \dots & x_{nk} \end{bmatrix}$$

Principal
Components

Weight Vectors

Original Data

Demo

**Implement principal components
analysis in scikit-learn**

Summary

Clustering is an unsupervised learning technique which helps find patterns in data

Common clustering algorithms are k-means, mean-shift clustering

Dimensionality reduction represents inputs in terms of their most significant features


PCA is a very commonly used technique for latent factor analysis



Books

**Hands-On Machine Learning with
Scikit-Learn and TensorFlow**

by Aurélien Géron



Related Courses

How to Think About Machine Learning Algorithms

Understanding Machine Learning with Python

Understanding the Foundations of TensorFlow